3/3

Hub

3 Lessons 45min

Next up: Discovering the Orchestration Service

Try Again

Continue

SAP Learning

Quick links

nload Catalog (CSV, JSON, XLSX, XML)

earning Hub

Training Shop

Developer Center

Community

letter

ng Support

upport

e Feedback

se Notes

SAP

pany Information

ers

s and Press

ormation

cy

s of Use

Disclosure

ot Share/Sell My Personal Information (US Learners Only)

rências de Cookies

Discovering Generative AI Hub Interface

Objective

After completing this lesson, you will be able to navigate the generative AI hub interface.

Discovering Generative AI Hub Interface

Once you have accessed the generative AI hub and deployed models, the next step is to get familiar with the user interface of SAP AI Launchpad, which is used to access the generative AI hub. Using this environment is not only a great way to explore multiple models and enhance your day-to-day productivity, but it also provides robust features to manage your prompts and LLM application lifecycles, ultimately helping you address complex business problems more effectively.

This lesson will provide you with an overview of the SAP AI Launchpad, with a focus on the generative AI hub. You will see how the features of the generative AI hub provide a productive environment for using multiple models, managing your prompts, and orchestrating workflows all in one place in a secure and reliable manner.

SAP AI Launchpad Interface

SAP AI Launchpad is a multi-tenant software as a service (SaaS) application on SAP BTP. You can use SAP AI Launchpad to manage AI use cases, known as scenarios, across multiple instances of AI run times, for example, SAP AI Core.

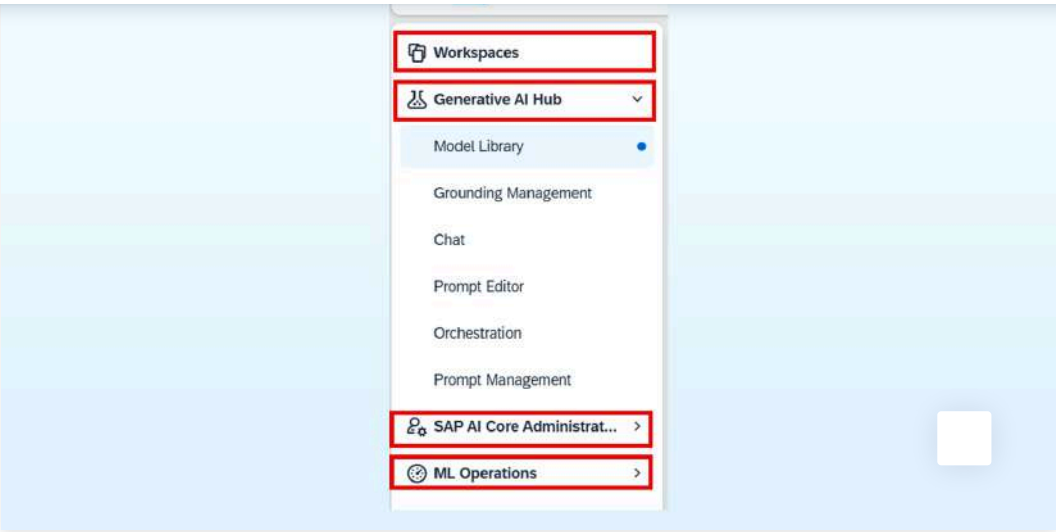

> **Note**
> These apps are only accessible to administrators and developers in your organization. As a business user, you will access and use generative AI hub described later.

Workspaces: This app allows technical users to manage API connections and organize resources. Each API connection represents a unique tenant, which can further contain multiple resource groups for streamlined generative AI hub access management.

Generative AI hub: This is the core application we will focus on. It provides all the necessary functionalities for working with generative AI, including managing models, developing prompts, and overseeing the lifecycle of your LLM applications.

SAP AI core Administrator: This app is used by administrators and developers to manage underlying infrastructure components such as Git repositories, application deployments, Docker registry secrets, and various connection secrets. It ensures secure and authorized access to the generative AI hub's capabilities.

ML Operations: This app is used to manage scenarios that can be used to create executables and configurations. These executable elements can be used to deploy models and orchestration service in the generative AI hub. This application also provides support for bringing in your datasets,

The generative AI hub within SAP AI Launchpad offers a suite of interfaces, including model library, chat, prompt editor, prompt management, orchestration, and administration.

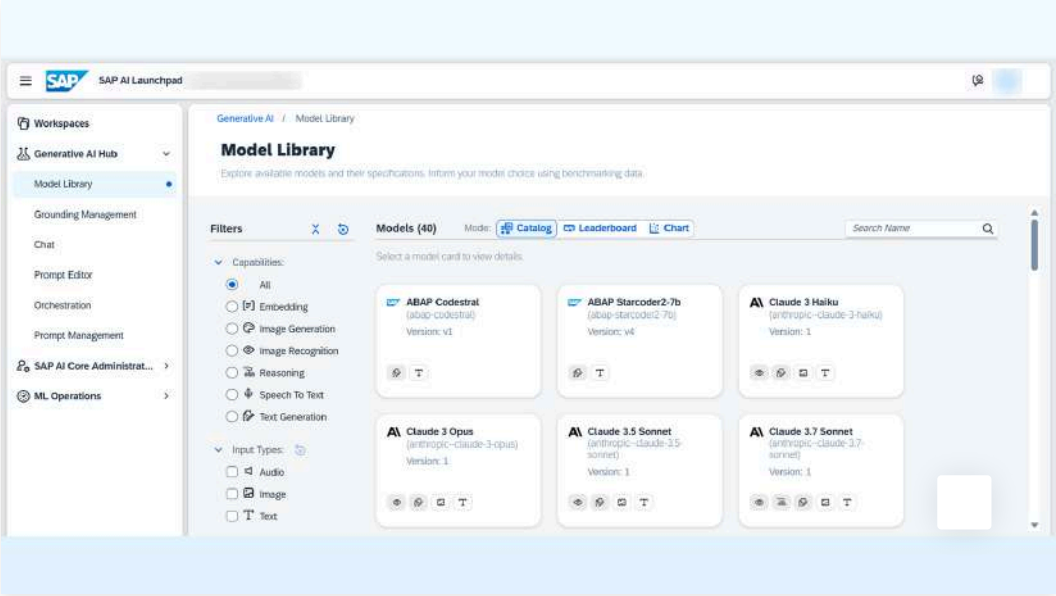

Model Library

The Model Library shows the list of all the available models in the generative AI hub and serves as a resource for selecting appropriate models. It provides comprehensive information to help you decide which model best fits your use case, considering factors like performance, cost, and specialization.

It provides the following modes that can be used to select a model.

Catalog Mode:

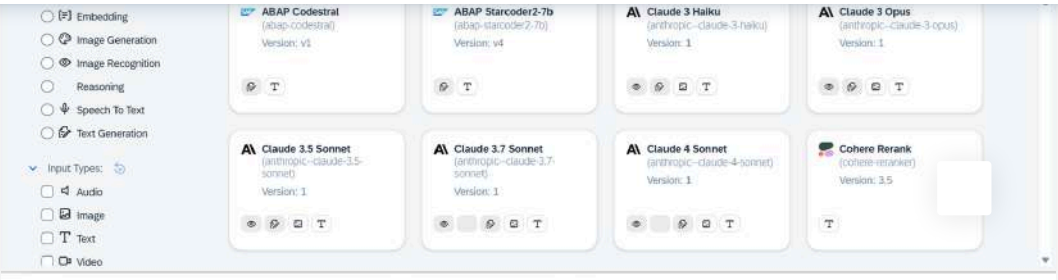

This mode shows all the available models with their metadata. Use filters to refine your selection or search for a model by name. A **model card** is displayed when you select a model.

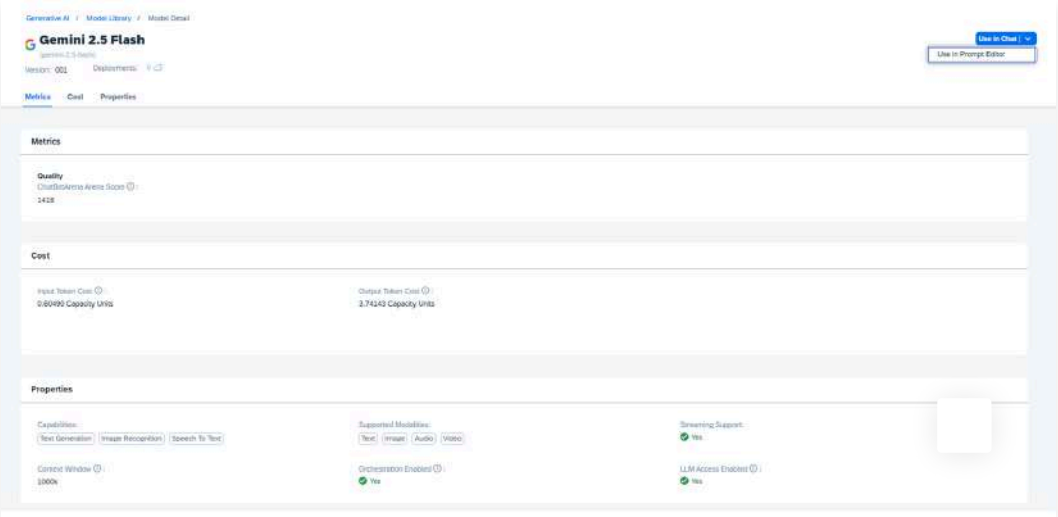

The Model Card offers detailed information, including data input types, potential costs, and relevant performance metrics.

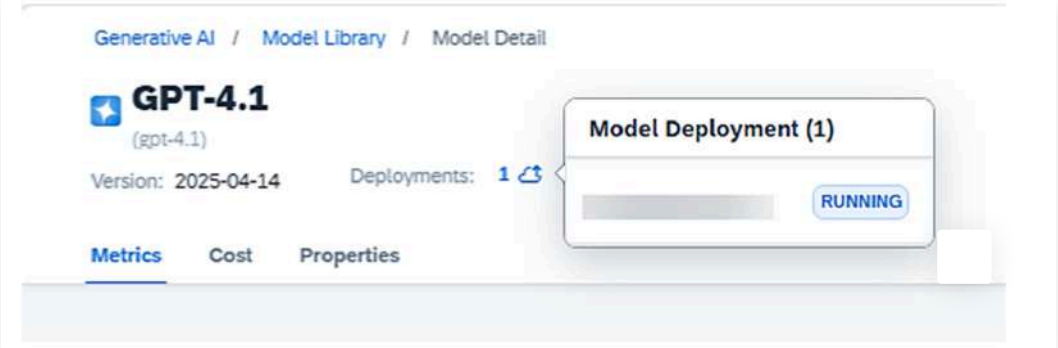

This is where you can also initiate deployment or check the status of a deployed model.

Leaderboard Mode:

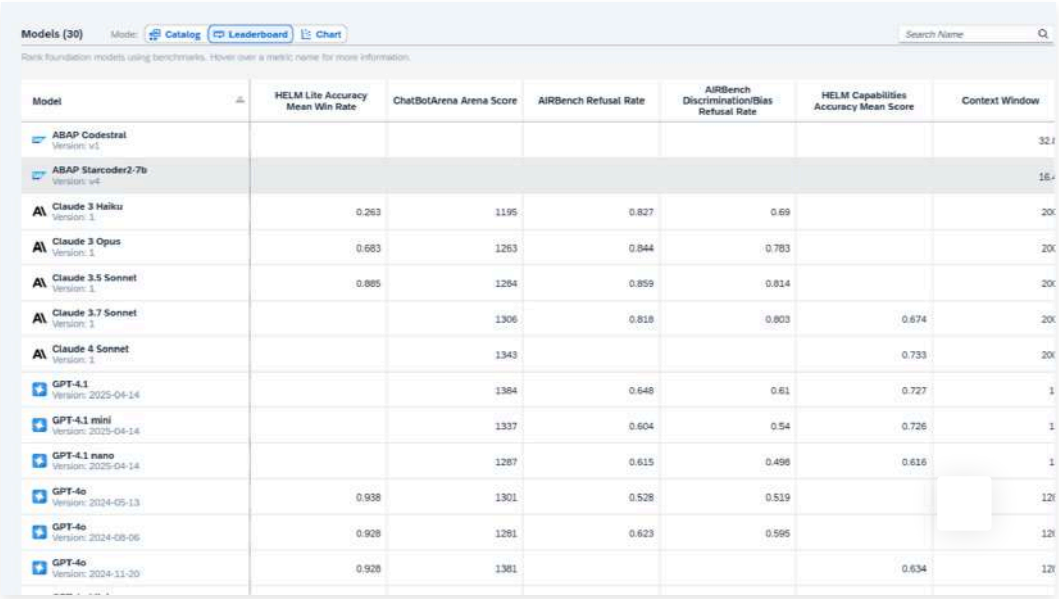

Models (30) Mode: Catalog Leaderboard Chart Search Name

Rank foundation models using benchmarks. Hover over a metric name for more information.

Model	HELM Lite Accuracy Mean Win Rate	ChatBotArena Arena Score	AIRBench Refusal Rate	AIRBench Discrimination/Bias Refusal Rate	HELM Capabilities Accuracy Mean Score	Context Window
ABAP Codestral Version: v1						32.k
ABAP Starcoder2-7b Version: v4						16.k
Claude 3 Haiku Version: 1	0.263	1195	0.827	0.69		20
Claude 3 Opus Version: 1	0.683	1253	0.844	0.783		20
Claude 3.5 Sonnet Version: 1	0.885	1284	0.859	0.814		20
Claude 3.7 Sonnet Version: 1		1306	0.818	0.803	0.674	20
Claude 4 Sonnet Version: 1		1343			0.733	20
GPT-4.1 Version: 2025-04-14		1384	0.648	0.61	0.727	1
GPT-4.1 mini Version: 2025-04-14		1337	0.604	0.54	0.726	1
GPT-4.1 nano Version: 2025-04-14		1287	0.615	0.498	0.616	1
GPT-4o Version: 2024-05-13	0.938	1301	0.528	0.519		12
GPT-4o Version: 2024-08-06	0.928	1281	0.623	0.595		12
GPT-4o Version: 2024-11-20	0.928	1381			0.634	12

This mode shows model scores across a range of benchmarks. You can apply filters to narrow down your options, search by name, or reorder the list based on specific benchmarks.

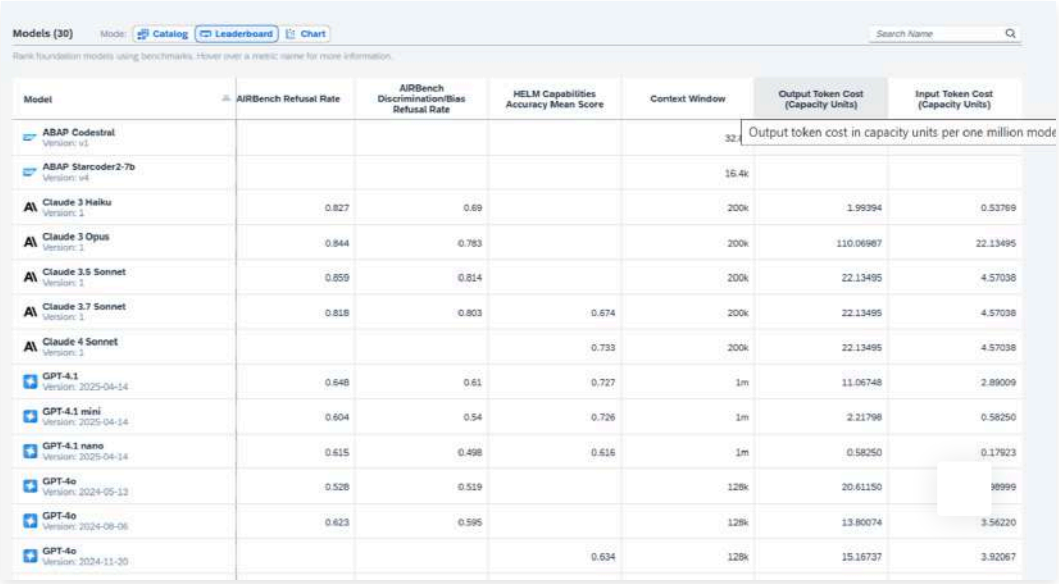

Models (30) Mode: Catalog Leaderboard Chart Search Name

Rank foundation models using benchmarks. Hover over a metric name for more information.

Model	AIRBench Refusal Rate	AIRBench Discrimination/Bias Refusal Rate	HELM Capabilities Accuracy Mean Score	Context Window	Output Token Cost (Capacity Units)	Input Token Cost (Capacity Units)
ABAP Codestral Version: v1				32.k	Output token cost in capacity units per one million mode	
ABAP Starcoder2-7b Version: v4				16.4k		
Claude 3 Haiku Version: 1	0.827	0.69		200k	1.99394	0.53769
Claude 3 Opus Version: 1	0.844	0.783		200k	110.06987	22.13495
Claude 3.5 Sonnet Version: 1	0.859	0.814		200k	22.13495	4.57038
Claude 3.7 Sonnet Version: 1	0.818	0.803	0.674	200k	22.13495	4.57038
Claude 4 Sonnet Version: 1			0.733	200k	22.13495	4.57038
GPT-4.1 Version: 2025-04-14	0.648	0.61	0.727	1m	11.06748	2.89009
GPT-4.1 mini Version: 2025-04-14	0.604	0.54	0.726	1m	2.21798	0.58250
GPT-4.1 nano Version: 2025-04-14	0.615	0.498	0.616	1m	0.58250	0.17923
GPT-4o Version: 2024-05-13	0.528	0.519		128k	20.61150	8999
GPT-4o Version: 2024-08-06	0.623	0.595		128k	13.80074	3.56220
GPT-4o Version: 2024-11-20			0.634	128k	15.16737	3.92067

You also have the option of comparing context window and costs associated across multiple models and selecting the best model for your use case. Just hover over each benchmark or parameter to understand its

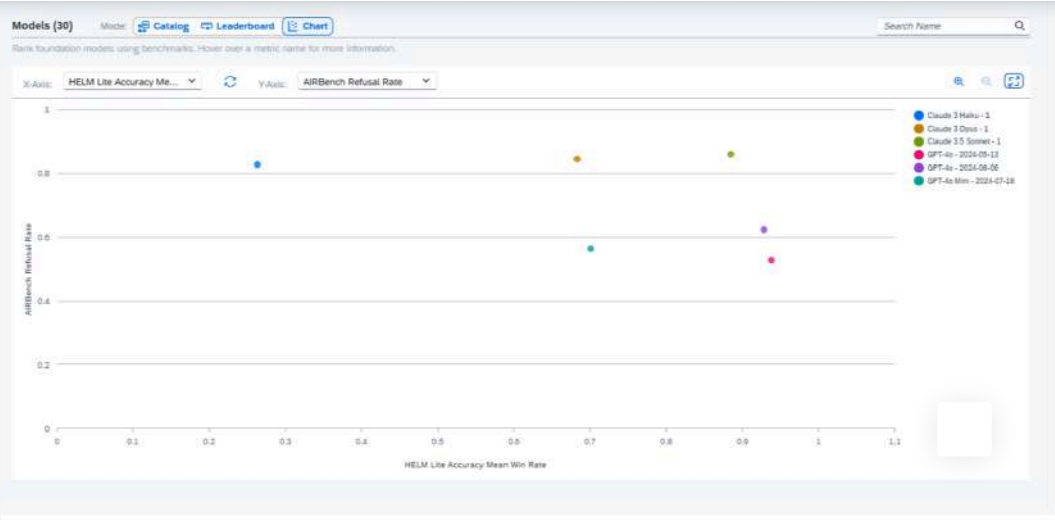

This mode provides a visual comparison between any two benchmarks or parameters.

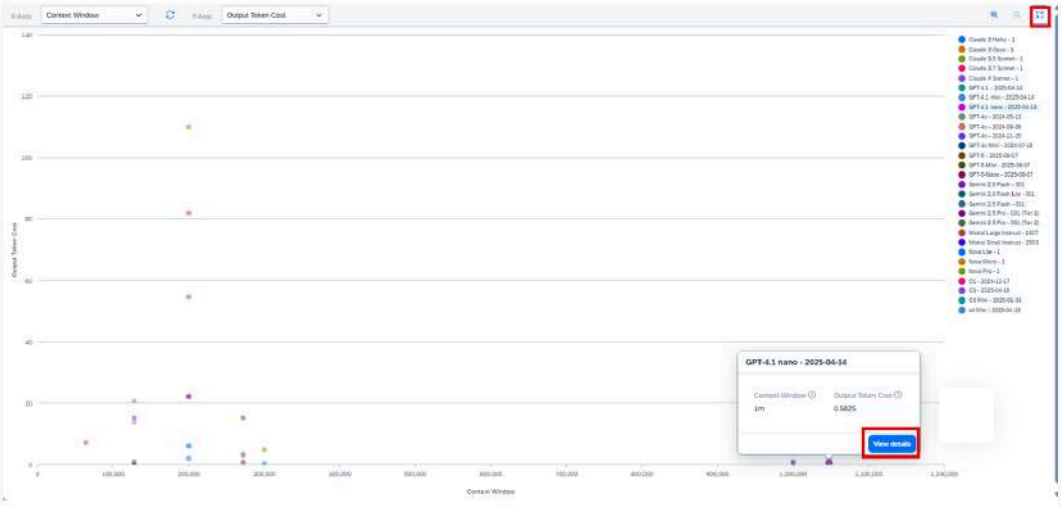

For example, you can choose a context window and output token cost and select a model with the largest context window and the lowest output token cost. You can see the model card in detail and select the optimum model for your use case. You can also zoom out a chart for better analysis. This is one of the many valuable features of generative AI hub to help you optimize your LLM applications and solve business problems on a scale.

For more information, visit the SAP AI Launchpad Help Portal - Model Library.

your solution is both efficient and budget friendly.

Chat Interface

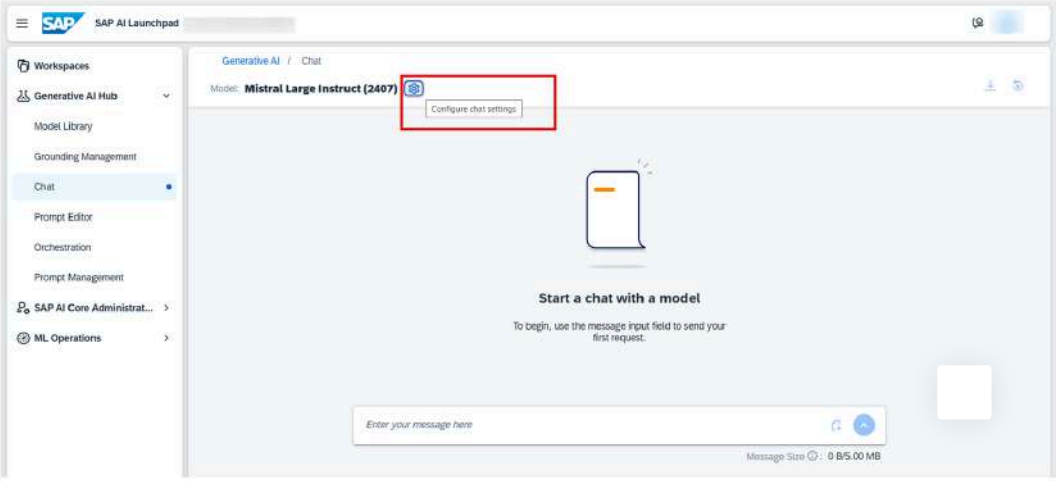

The chat interface serves as a conversational platform for interacting with models, enhancing user experience through personalized interactions. You can engage with models conversationally, with the context of previous messages maintained within the session. This feature enables developers to prototype and build applications that can scale efficiently.

You can select any available model using the Configure chat settings gear.

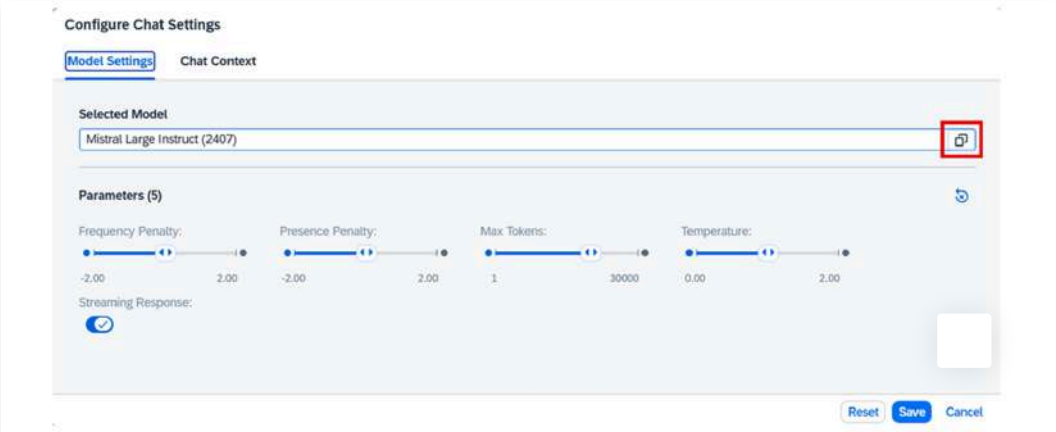

Use the select option to open the Model Library and choose the appropriate model for your case.

You can also update the available parameters for each model.

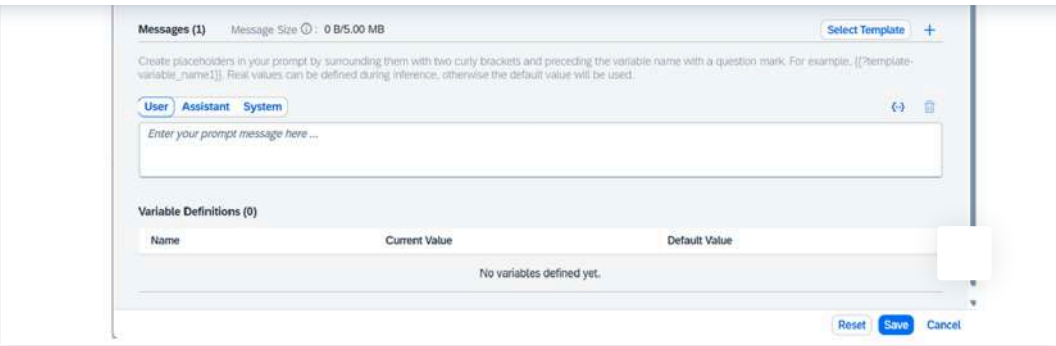

You can also update the context history for a larger context in the chat conversation and use advanced features like message roles and variables to tailor the conversation for your business needs.

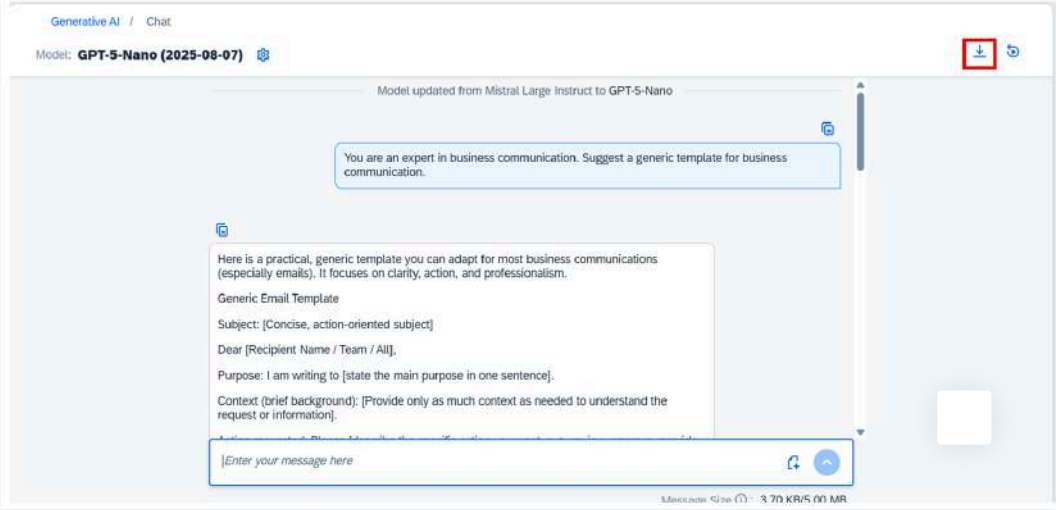

You may download your chat by selecting the **Download** icon. Text data will be saved locally in the JSON format; note that images are excluded from downloads.

To copy text data from an individual chat message or response, use the **Copy** icon. Images cannot be copied.

Use Case Example:

You have an idea for an AI assistant that helps sales teams answer common product questions. You can use the Chat interface to quickly test different LLMs with sample questions ("What are the key features of S/4HANA Cloud?" or "How does our product compare to a competitor's?"). This allows you to rapidly prototype conversational flows and see how well various models understand and respond before building out a full

The grounding management interface is used to manage the lifecycle of data pipelines. These data pipelines are used to provide more context for LLMs using documents stored in a SharePoint or cloud storage like S3.

For more information about grounding techniques, visit the Using Advanced AI Techniques with SAP's Generative AI Hub course.

Prompt Editor

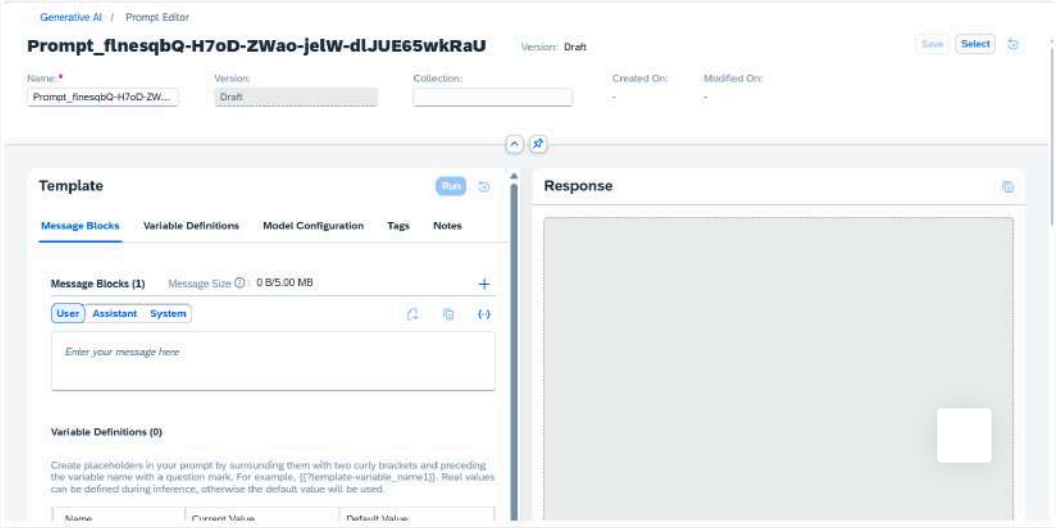

The prompt editor provides a comprehensive environment for creating prompts, selecting models, saving prompts, and configuring parameters.

Key Features:

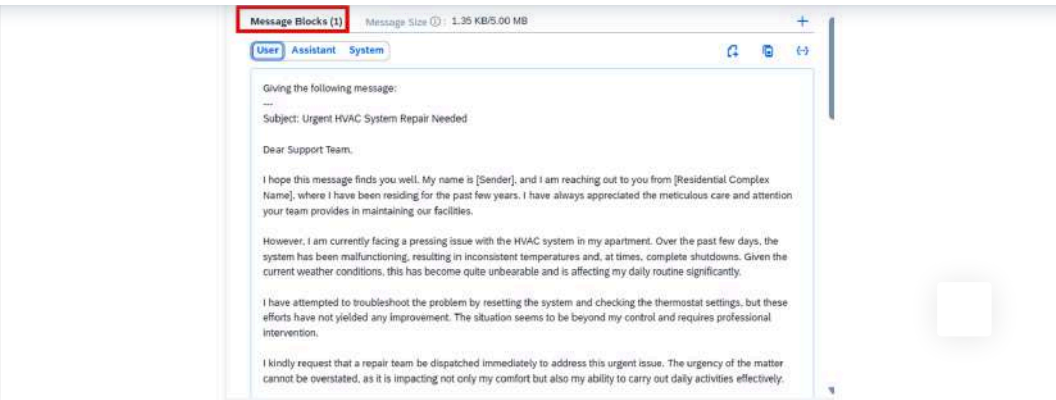

Message Blocks: This is a dedicated area for writing prompts for the System, Assistant, and User roles. You can add multiple blocks for multiple roles in a version.

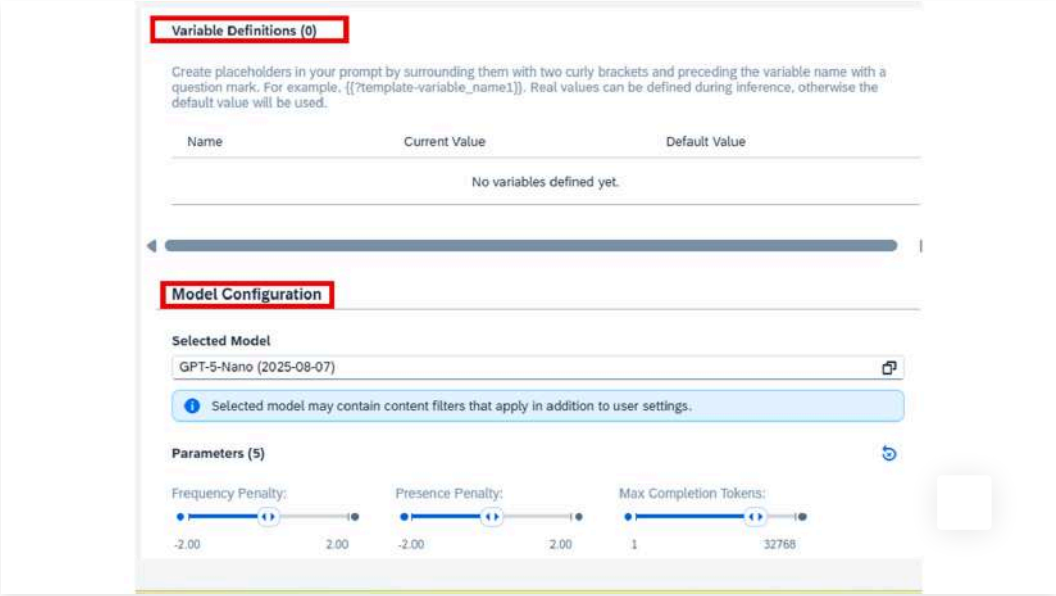

Variable Definition: You can add variables to your prompts to make them flexible and customizable in various situations. You can reuse a prompt for your use case by using default values or current values. Using the SDK, variables let you input information at runtime, either through user interfaces or system workflows.

Model Selection: You can select from various models available in the generative AI hub and modify their parameters directly within the editor to achieve the best results tailored to your specific use cases.

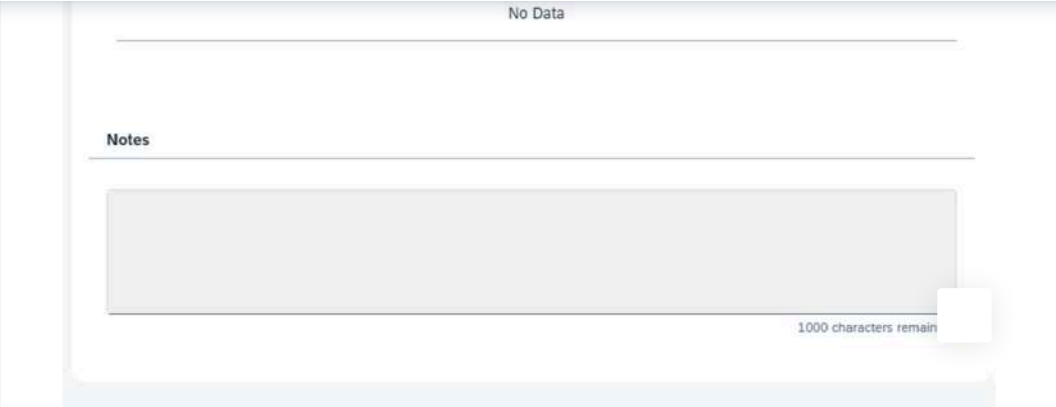

Metadata: You can add organizational metadata, such as tags and notes, to categorize and manage your prompts effectively.

Save Prompt and Version Management: The Prompt Editor allows you to save prompts and even assign a collection name to organize them. This ensures traceability and allows for iterative refinement.

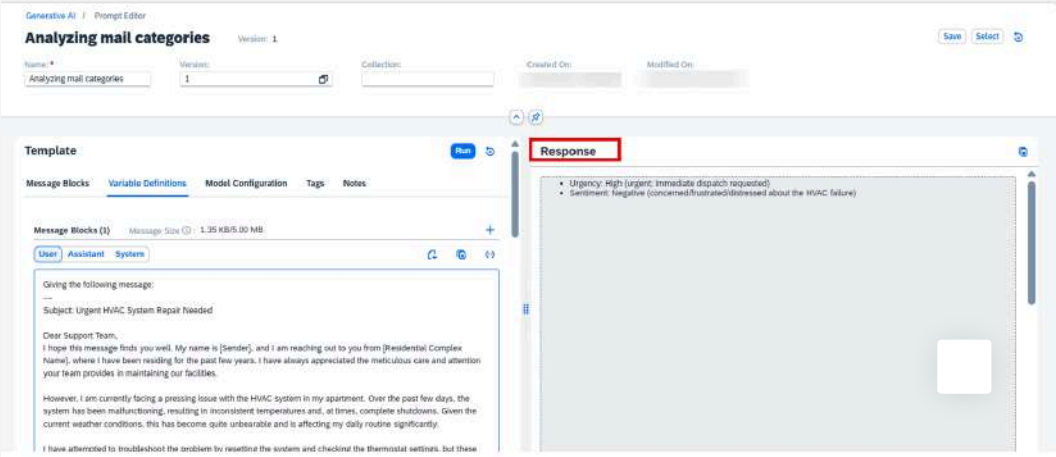

Response: You can see the response and save the results. You can continue editing the prompt in the same window and save multiple versions for future reference.

Use Case Example: Your business needs to generate personalized marketing emails for different customer segments. Using the Prompt Editor,

This interface is the go-to place for creating, managing, and retrieving prompts for prototyping and prompt lifecycle management.

This interface in the AI hub also supports human-to-AI and software-to-AI interactions. You can use direct prompts or feed software-generated reports for further analysis.

Prompt Management

Once you save a prompt, you can access and reuse it in Prompt Management.

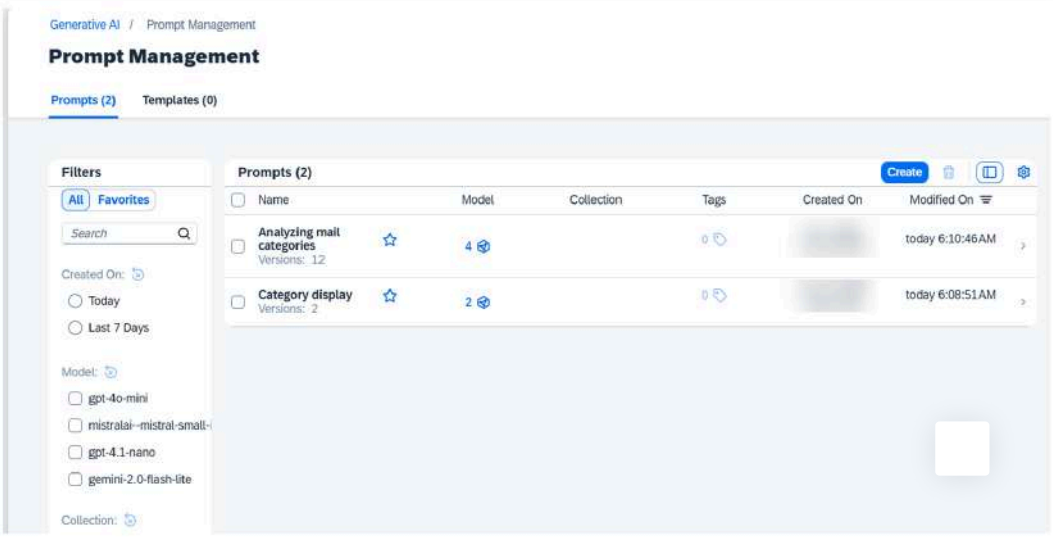

You can see all the versions and models used in the prompt.

Learning

Quick links

Download Catalog (CSV, JSON, XLSX, XML)

SAP Learning Hub

SAP Training Shop

SAP Developer Center

SAP Community

Newsletter

Learning Support

Get Support

Share Feedback

Release Notes

About SAP

Company Information

Copyright

Trademark

Worldwide Directory

Careers

News and Press

Site Information

Privacy

Terms of Use

Legal Disclosure

Do Not Share/Sell My Personal Information (US Learners Only)

Preferências de Cookies

Explaining Generative AI Hub Applications

Objective

After completing this lesson, you will be able to explain generative AI applications beyond chatbots, using software-driven input and output processing scenarios.

Explaining Generative AI Hub Applications

This lesson will broaden your perspective on the versatility of Generative AI. We will delve into a range of practical applications that leverage the SAP Generative AI Hub to solve complex business problems, automate tasks, enhance decision-making, and drive innovation across the enterprise. You will discover how Generative AI can interact with human users and software systems in sophisticated ways, unlocking unprecedented value.

Applications of Generative AI

Generative AI, particularly Foundational Models or FMs, has opened a wide range of potential applications beyond the traditional chatbot paradigm. Many still see generative AI as simply processing human input for output, but its potential goes much further.

Generative AI can independently produce reports, explanations, and precise translations without constant human input. This enables task automation and improves accessibility of information.

Generative AI can also interpret human instructions to control software, enabling intelligent interfaces that efficiently execute commands without

distinct, as many applications include overlapping components, and prompt engineering can integrate human input with software processes. Generative AI impacts numerous industries and domains, with uses extending beyond chatbots to support innovation and automation.

To further illustrate these four interaction types, let's look at specific examples within the SAP landscape:

- **Human-to-AI:** A business analyst uses a chat interface to ask, "Summarize last quarter's customer support tickets by product category," the AI provides a structured textual summary derived from the ticket data.
- **Software-to-AI:** An SAP backend system feeds sales order data to the AI, generating a natural language report highlighting fulfillment bottlenecks and regional trends.
- **Human-to-Software via AI:** A developer describes a desired feature, "Generate a boilerplate for an OData service in SAP BTP to manage product inventory," and the AI produces the corresponding code structure for the service.
- **Software-to-Software via AI:** An automated process in SAP SuccessFactors identifies an understaffed department, and an AI-driven service then initiates a recruitment workflow in a connected applicant tracking system.

Use Case for Generative AI hub

Example: Consistent Product Description Generation

Beyond email automation, consider a scenario where a manufacturing company struggles to maintain consistent, high-quality product descriptions across its numerous sales channels and marketing materials. Inconsistent descriptions lead to customer confusion and can erode brand trust.

By using the generative AI hub, the company can deploy a solution to create a unified product description generator. An LLM, provided with core product specifications, for example from an SAP Product Lifecycle Management system, can automatically generate detailed, engaging, and consistent product descriptions.

Value Proposition

This application ensures brand consistency across all platforms, significantly improves customer understanding of products, and can boost sales. Critically, it automates a tedious manual task, freeing up marketing and product teams to focus on more strategic initiatives. The generative AI hub provides secure access to the advanced AI models needed to understand product nuances and generate high-quality content.

The example illustrates that the generative AI hub is not just for creating conversational interfaces but is a versatile platform for building intelligent solutions that automate, analyze, and enhance processes across the enterprise by managing various forms of software-driven input and output.

Lesson Summary

automation and consistent product description generation, by enabling flexible model experimentation and seamless integration. This broader view of generative AI's capabilities provides a strong foundation and motivation for exploring the capabilities of the generative AI hub.

Next lesson

Was this lesson helpful? Yes No

Learning

Quick links

Download Catalog (CSV, JSON, XLSX, XML)

SAP Learning Hub

SAP Training Shop

SAP Developer Center

SAP Community

Newsletter

Learning Support

Get Support

Share Feedback

Release Notes

About SAP

Company Information

Copyright

Trademark

Worldwide Directory

Terms of Use

Legal Disclosure

Do Not Share/Sell My Personal Information (US Learners Only)

Preferências de Cookies

SAP Note

Language: English | Deutsch (?iv_language=D)

3237166 - AI Core Component Upgrade: Kserve 0.7 (/sap/support/notes/3237166) ☆

Component: CA-ML-AIC (SAP AI Core), Version: 1, Released On: 26.10.2022

Symptom

SAP AI Core Component Kserve upgraded to v0.7

Other Terms

AIW, AI Workbench, AI Core, AI

Reason and Prerequisites

SAP AI Core Component Kserve upgraded to v0.7 for security compliance

Solution

Changes:

In ServingTemplate:

1. The ***spec.template.apiVersion*** to be changed from **serving.kubeflow.org/v1beta1** to **serving.kserve.io/v1beta1**
2. The ***spec.template.spec.predictor.containers.name*** to be changed from **kfserving-container** to **kserve-container**
3. With kserve v0.7, there's no support for apiVersion: **serving.kubeflow.org/v1alpha2**

Updates:

- All the **RUNNING deployments** will be migrated to support the new kserve version 0.7 by **AI Core**.
- All **ServingTemplate**'s must be migrated to support **serving.kserve.io/v1beta1** spec by **Users**. (See - User Action Required)

User Action Required:

- AI Core users are requested to update **ServingTemplate** to support **serving.kserve.io/v1beta1** spec by **September 30**.
- AI Core users, based on the current **spec.template.apiVersion**, can update the **ServingTemplate** manually, or use the Template Patch Tool (See - Resources)

Type	API Version	Template Update Required

Kfserving	*serving.kubeflow.org/v1beta1*	**No**
Kfserving	*serving.kubeflow.org/v1alpha2*	**Yes**
Kfserving	Not Specified (defaults to *serving.kubeflow.org/v1alpha2*)	**Yes**

Resources:

- Template Patch Tool: (**https://github.wdf.sap.corp/I343697/kserve-upgrade/tree/main/releases/2208a/patch_serving_template (https://github.wdf.sap.corp/I343697/kserve-upgrade/tree/main/releases/2208a/patch_serving_template)**)

This document is referenced by

SAP Note/KBA	Title
3248365	Central SAP Note for SAP AI Core (/sap/support/notes/3248365)

Legal

Privacy (http://www.sap.com/corporate-en/about/legal/privacy.html) |
Terms of use (https://support.sap.com/support-programs-services/about/terms-of-use.html) |
Legal Disclosure (http://www.sap.com/corporate-en/about/legal/impressum.html) |
Copyright (http://www.sap.com/corporate-en/about/legal/copyright/index.html) |
Trademark (http://www.sap.com/corporate-en/about/legal/copyright/index.html#trademark)

Follow

(https://www.facebook.com/SAPDigitalBusinessServices) (https://twitter.com/SAPSupportHelp)
(https://www.youtube.com/user/SAPSupportInfo) (https://www.linkedin.com/groups/138840)
(https://plus.google.com/+SAPCloud)

SAP ONE Support Launchpad

Home (https://launchpad.support.sap.com) | My Favorite Notes (https://launchpad.support.sap.com/#/mynotes?tab=Favorites) | Expert Search (https://launchpad.support.sap.com/#/mynotes?tab=Search)

3248365 - Nota SAP central para SAP AI Core

Componente: CA-ML-AIC (Componentes válidos para várias aplicações > Machine Learning > SAP AI Core), Versão: 4, Liberado em: 22.07.2025

⚠ Atenção Este é um documento traduzido automaticamente que não foi revisado para garantir acuracidade.

Sintoma

Esta nota SAP fornece informações gerais sobre o SAP AI Core.

Solução

Recursos e ampliações

Para saber mais sobre as novas capacidades e ampliações introduzidas no SAP AI Core, consulte o documento O que há de novo no SAP Help Portal.

Assistência ao usuário

O conteúdo de assistência ao usuário é publicado no SAP Help Portal.

Este documento se refere a

Nota SAP/KBA	Componente	Título
3634540		CA-ML-AIC
3505347	CA-ML-AIC	Orquestração
3491182	CA-ML-AIC	Implementações do centro de produção e disponibilidade do plano de recursos
3437766	CA-ML-AIC	Disponibilidade de modelos de IA generativa
3431716	CA-ML-AIC	SAP AI Core: Implementação de Alerta de Segurança e Hotfix
3239609	CA-ML-AIC	Desativação de pontos de extremidade de API obsoletos na API do AI Core
3237166	CA-ML-AIC	Upgrade de componente do AI Core: Kserve 0.7

Exploring the Generative AI Hub

Objective

After completing this lesson, you will be able to identify the AI foundation and generative AI hub.

Exploring Generative AI Hub

As artificial intelligence rapidly reshapes business landscapes, organizations seek powerful yet secure and integrated ways to harness its potential. For many, that journey begins with understanding the core infrastructure that enables AI-driven innovation.

As a leader in enterprise applications, SAP is committed to infusing intelligent capabilities directly into the core of your business processes through SAP Business AI.

This lesson will introduce you to the fundamental infrastructure that makes this possible: SAP's AI Foundation. You will explore a key component of this foundation, the generative AI hub. You will see its role as your gateway to securely and effectively leverage Large Language Models (LLMs) for enterprise-grade solutions. Whether you're a business leader, a consultant, or a technical expert, understanding this layered approach is crucial for unlocking the full potential of generative AI within your organization.

AI Foundation

AI is a strategic need for organizations. However, as they seek to innovate faster and operate more efficiently, they often face growing complexity in managing AI development, deployment, and governance. This is precisely where SAP's AI Foundation comes in.

The AI Foundation serves as **SAP's AI operating system,** enabling organizations to **build, extend, and operate** custom AI solutions at scale in a secure, responsible, and efficient manner on SAP Business Technology Platform (SAP BTP).

It is the backbone of SAP Business AI, providing the infrastructure, tools, and services needed to orchestrate cutting-edge technology with deep business context.

AI Foundation Matters:

Organizations today face several challenges, including high resource demands for onboarding AI models, fragmented tools, and difficulty scaling AI responsibly across regions and regulations. AI Foundation directly addresses these by offering:

- **A Unified Development and Orchestration Platform:** Providing a central, integrated access point for all AI Foundation tools, simplifying collaboration and governance.
- **Built-in Governance and Security:** Ensuring ethical, secure, and compliant AI operations from the ground up, significantly reducing risk.
- **Seamless Integration:** Connecting effortlessly with SAP and non-SAP data, infrastructure, and partner models, allowing for broad applicability and leveraging existing investments.

Core Products and Architecture of AI Foundation

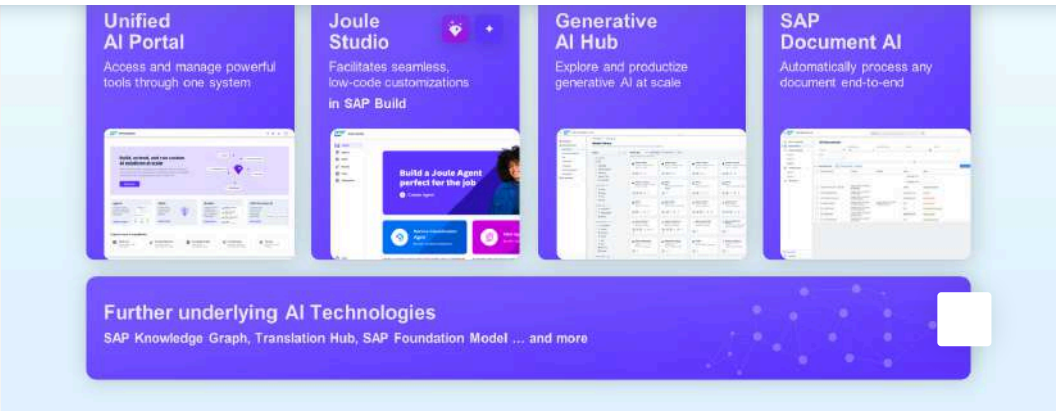

SAP's AI Foundation is built upon a sophisticated, multi-layered architecture designed to support enterprise-grade AI operations and deliver on SAP's core AI promises. It combines well-established components and makes them available through a single entry point: the Unified AI Portal. This portal serves as a centralized access hub for all AI Foundation tools, simplifying collaboration and lifecycle management.

Key products integrated into AI Foundation include:

- **Joule Studio:** A low-code/no-code environment specifically for building and customizing AI agents and extending Joule skills, empowering developers and business users to innovate quickly.
- **Generative AI Hub:** This crucial component provides secure and managed access to top-tier LLMs and advanced prompt engineering tools, a core focus of this course. It accelerates AI development by enabling controlled and efficient LLM utilization.
- **SAP Document AI:** Automates data handling across a wide range of unstructured documents, significantly enhancing the accuracy and quality of data processing within critical business workflows.
- **Underlying Technologies:**
  - **SAP Knowledge Graph:** Connects vast amounts of SAP and external data, creating a rich web of contextual intelligence that enables the development of smarter, more relevant AI outcomes.
  - **SAP Foundation Model:** Produces reliable predictions based on SAP application data and addresses some limitations of AI and LLMs.

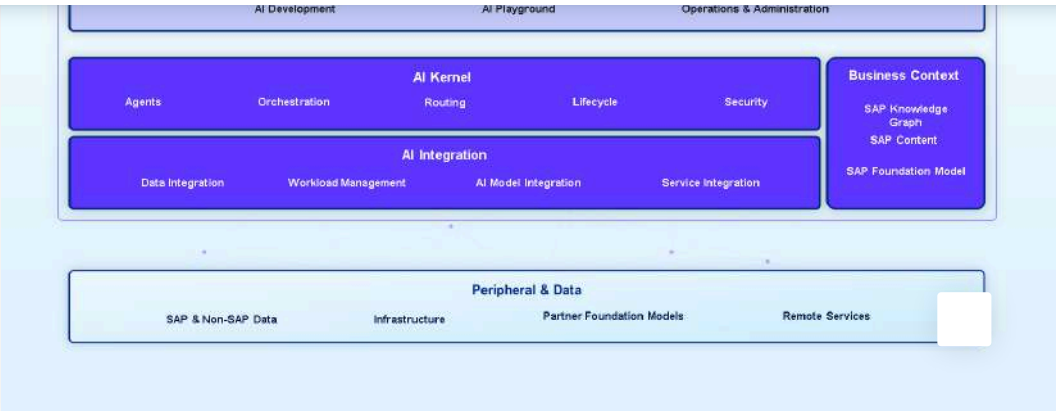

The AI Foundation's architecture is structured across four key layers to ensure robust and scalable AI operations:

1. **OS Interfaces Layer:** This serves as the unified entry point for development, operations, and administration. It is where user-facing tools like the AI Playground and Joule Studio reside, enabling intuitive interaction.
2. **AI Kernel Layer:** This core layer manages AI agents, workloads, models, and their entire lifecycle. It handles critical tasks like resource allocation and scheduling to ensure secure, compliant, and efficient AI operations.
3. **AI Integration Layer:** This layer facilitates seamless connectivity to diverse data sources, various models, and various services. It supports workload management, data integration pipelines, and the sophisticated orchestration of AI models across different environments.
4. **Peripheral and Data Layer:** This foundational layer provides the essential infrastructure by integrating both SAP and non-SAP data sources, leveraging existing IT infrastructure, and connecting to partner models. It also supplies the raw materials and computational power needed for AI processes.

This layered structure ensures that AI Foundation is not just a collection of tools, but a cohesive AI operating system that serves as the backbone for all AI technologies and solutions that SAP internally utilizes. Simultaneously, it empowers customers and developers with a comprehensive suite of tools to build, extend, and run their custom AI solutions and agents on a scale. SAP also fosters a strong partner ecosystem through its open approach, allowing seamless integration with leading AI and data partners to leverage best-in-class models and industry expertise.

Resource Groups

SAP AI Core uses resource groups to keep different machine learning resources and tasks separate.

A resource group is a virtual folder for related items inside one SAP AI Core account. When you get started, you automatically get a default resource group. You can add or remove more resource groups if you're an administrator using SAP AI Launchpad. Resource groups may be configured according to the specific requirements of individual tenants and their respective use cases.

The Generative AI Hub: Powering LLMs within the AI Foundation

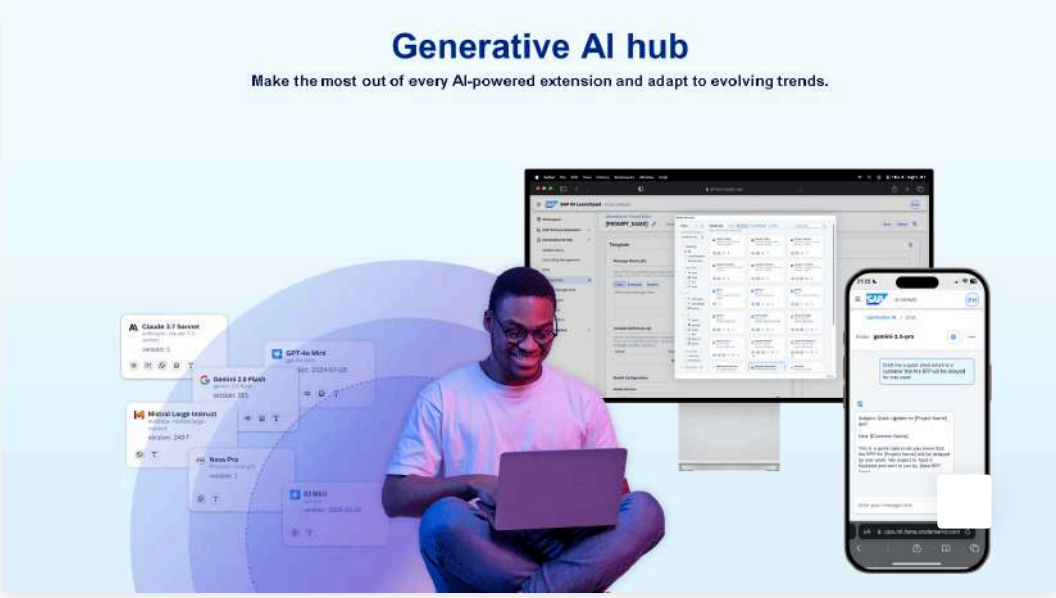

The growing landscape of LLMs offers unprecedented potential, but also introduces complexities in systematic model selection, access, and governance. The **generative AI hub** is an essential part of SAP's AI Foundation, within SAP AI Core, and is designed to meet enterprise requirements. It is designed to enable systematic and tool-assisted model selection, access, and exploration of LLMs, ensuring they are based on business and contextual data.

The generative AI hub is fundamental to SAP's promise of **Relevant, Reliable, and Responsible AI** when leveraging LLMs:

possible through features like the Grounding powered by SAP HANA Vector Engine. This makes it easier to find business documents related to a query and use them as context for the LLM.

- **Reliable (Trusted Access and Performance):** It consolidates trusted LLM and foundation model access, reducing dependency on a single vendor by supporting diverse models from multiple providers. This ensures the continuous functionality of your AI applications, with features like automatic and manual upgrade options for model versions.
- **Responsible (Security and Compliance by Design):** Generative AI hub ensures compliance and trustworthiness through built-in governance, security measures, and ethical considerations. It includes critical functionalities like content filtering and data masking to prevent risks such as prompt injection and jailbreaking, as well as adhering to strict security standards like SOC 2, NIST, and ISO certifications. Your data is fully protected and isolated in the generative AI hub. SAP does not share customer data with external vendors for model training. Customer data is never accessible to other customers. SAP uses enterprise-grade LLMs that do not store API input or output, ensuring complete protection of customer information.

Key Capabilities of the Generative AI Hub

The Generative AI Hub covers the entire lifecycle of Generative AI applications:

- **Access to multiple models:** Provides streamlined access to frontier AI models from various providers and flexible compute resource allocation. This ensures you can easily connect to the best LLMs for your specific use cases.
- **Model agnostic development:** You can easily switch between models, adapt orchestration configurations on the fly, and manage prompt variants for different LLMs. This capability means you can drive AI adaptation at your pace, leveraging new and improved models without disrupting your existing solutions or requiring migrations to other services.
- **Exploration and Development:** Offers a comprehensive toolset for building custom AI solutions. This includes a prompt editor, robust prompt management (registry, libraries), and Software Development

robust SDKs and client libraries, to efficiently integrate generative AI into your applications and accelerate the development of bespoke AI solutions. SAP Cloud SDK for AI is the official SDK for SAP AI Core, the generative AI hub, and orchestration. It provides model access, enabling the integration of the generative AI hub capabilities with common frameworks like SpringAI, Langchain, and Langgraph.

- **Orchestration:** Facilitates coordinating and managing various AI components, including LLMs. It enables the execution of multiple AI models, manages data flow, optimizes computational resources, integrates features like content moderation and data masking, and streamlines end-to-end AI workflows.
- **Deployment and Delivery:** Software deployment involves the steps, processes, and activities necessary to make a software system or update available to its intended users. Currently, most IT organizations and software developers use manual and automated methods to deploy software updates, patches, and new applications. You bring your own models (for example, fine-tuned open-source models, proprietary models) and deploy them. These capabilities enable the efficient deployment and delivery of AI models and applications, ensuring that they are operational and accessible to users.
- **Governance:** Implements robust policies and procedures to manage AI development, deployment, and operation in compliance with regulatory and organizational standards. Features include comprehensive logging for auditing, metering, monitoring, multi-tenancy support, and clear roles and responsibilities.

Support for the Model Lifecycle

The generative AI hub supports different approaches for using AI models:

Self-Hosted Models: Organizations can deploy their own models (such as fine-tuned open-source or proprietary models) and manage model artifacts, dependencies, and the entire lifecycle. This includes overseeing infrastructure, scaling, security, compliance, monitoring, costs, updates, versioning, and performance.

SAP Managed Models: This approach provides pre-integrated, managed access to various LLMs from third-party providers (such as Azure OpenAI GPT-5, Google Gemini 2.5 Pro, AWS Claude Sonnet 4, and others). Users do not need to obtain separate licenses for these LLMs, as SAP manages commercial agreements and technical integration. SAP is also responsible

retire older ones, and may adjust aspects such as performance and cost. As a result, the specific LLM version an application uses may not always remain available.

Generative AI hub ensures that these LLM version changes do not impact your applications by supporting auto upgrade and manual upgrades for the latest version of a particular LLM.. You can monitor model version deprecation dates, use either auto-upgrade or manual upgrade strategies to maintain continued service and benefit from the latest advancements.

SAP Hosted Models:

These models are hosted and managed by SAP in its own infrastructure, ensuring seamless integration with SAP systems and applications. SAP handles all aspects of security, compliance, performance optimization, and operational management, allowing organizations to leverage AI capabilities without the need to manage infrastructure or technical complexities. These models are pre-configured for use within the SAP ecosystem, providing reliable and secure AI-driven solutions tailored to business needs.

For all the latest information on all the available models in generative AI hub refer Note 3437766 .

Access Generative AI hub

To get started with the generative AI hub, users must establish access by setting up their SAP BTP global account, provisioning SAP AI Core, and ensuring the appropriate roles and permissions are assigned. The generative AI hub operates within SAP AI Core and is primarily accessed through the SAP AI Launchpad interface or programmatically via APIs and SDKs. With access secured, users can deploy LLMs from various leading providers, configuring essential parameters such as model type and version.

Each deployed LLM is provided with a secure endpoint for integrating enterprise applications, simplifying consumption and management with support for version monitoring and upgrades. This foundational process allows organizations to quickly harness generative AI for business innovation within their SAP environment.

Lesson Summary

You've now explored the foundational infrastructure that enables enterprise-grade AI within the SAP ecosystem. You understand how **SAP's**

generative AI applications are grounded in your unique business data, delivering trustworthy, context-aware, and high-value outcomes for your organization. It supports model lifecycle management for your LLM applications to ensure business continuity along with the latest model advancements.

Free COURSE • 1 hr +

Discovering SAP's Generative AI Hub

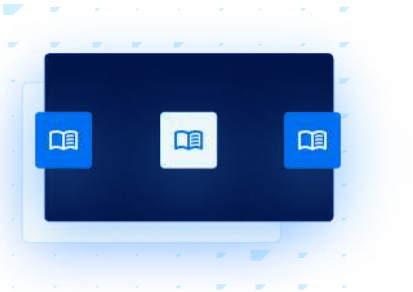

Start learning

Learning outcome

After completing this course,

- You will understand the architecture and main features of SAP's generative AI hub and its integration with SAP AI Core and AI Launchpad.
- You will be able to describe how the hub provides secure access to multiple language models, supports model selection and lifecycle management, and enables prompt engineering and...

See More

Course information

What you'll learn

UNIT 1

Describing the Generative AI Hub

3 Lessons 45 min

UNIT 2

Discovering the Orchestration Service

2 Lessons 35 min

Level up your skills

Subscription

Optional but recommended to boost your expertise.

LIVE SESSION

Generative AI Agents

LIVE SESSION

Build Your Own AI Solution with Generative AI hub

LIVE SESSION

Getting Started with SAP Generative AI Launchpad

LIVE SESSION

Get Certified: SAP Certified - SAP Generative AI Developer

SAP Learning

Quick links

- Download Catalog (CSV, JSON, XLSX, XML)
- SAP Learning Hub
- SAP Training Shop
- SAP Developer Center
- SAP Community
- Newsletter

Learning Support

Get Support

Share Feedback

Release Notes

About SAP

Company Information

Copyright

Trademark

Worldwide Directory

Careers

News and Press

Site Information

Privacy

Terms of Use

Legal Disclosure

Do Not Share/Sell My Personal Information (US Learners Only)

Preferências de Cookies